

Software Fault Isolation using the CompCert compiler

Alexandre Dang

Team Celtique

June 10, 2016

Flash vulnerable plugin

Do you know this logo?

Flash is famous for its multiple vulnerabilities

- consequences on Flash
- but ALSO endangers your browser



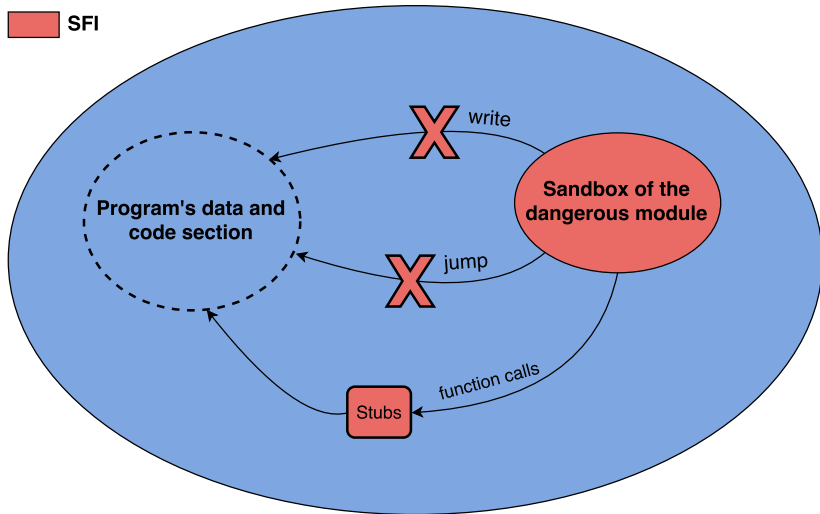
Goals of Software Fault Isolation (SFI)

- ▶ SFI aims to allow a protected program to execute dangerous modules in its own memory space without dangers.
- ▶ SFI confines the execution of the dangerous modules in a reserved area called sandbox
- ▶ `jump` and `write` instructions are protected by runtime checks
- ▶ function calls to the protected programs are controlled by SFI

Goals of SFI

 Memory of the protected program


 SFI



Overview of SFI

SFI chain is composed of two elements:


- ▶ the **generator** transforms the assembly code of the dangerous modules in order to confine the modules in their sandbox
- ▶ the **verifier** checks that the SFI transformations are present and valid before loading the code in memory



Sandboxing

Sandbox are continuous area identified by a tag.

For example the sandbox [0xda000000 - 0xdaffffff] has the tag 0xda:



Implementations

- ▶ First implementations for RISC architecture
- ▶ NativeClient, SFI for Google Chrome x86-32, x86-64 and ARM
 - ▶ Most complete implementation of SFI
- ▶ Portable Software Fault Isolation, implementation with the certified compiler CompCert
 - ▶ Take advantages of the correctness of CompCert
 - ▶ CompCert is the compiler used for our work

Pros and cons of SFI

- ▶ Pros

- ▶ Trusted Computing Base reduced to the verifier only
- ▶ Faster than protection by process separation

- ▶ Cons

- ▶ Architecture dependant
- ▶ Slows down the modified modules (between 5% and 21% depending on the implementation)

Problematics of SFI

- ▶ SFI has difficulties to deal with indirect jump through return addresses
- ▶ SFI is still vulnerable to Return Oriented Programming (ROP) attacks
- ▶ ROP attacks are one of the most common attacks in the industry
- ▶ We propose a solution to solve this issue

ROP attack example (1/3)

```
1 void evil_code() {  
2     printf("Argh, we got hacked!\n");  
3 }  
4  
5 void foo(char* input){  
6     char buf[1];  
7     ... code ...  
8     strcpy(buf, input);  
9     ... code ...  
10 }
```

ROP attack example (2/3)

```
terminal$ ./buffer $(python -c 'print 13*"a" +  
"\x7b\x84\x04\x08"')
```

Address of evil_code = 0x0804847b

Stack before:

0xf7712000

0xff957998

0xf7593d26

0xf7712d60

0x0804868c

0xff957978

0xf7593d00

0xf7713dc0

0xf77828f8

0xff957998

0x08048510

//Return address of *foo*

ROP attack example (3/3)

Stack after :

0xff958161

0xff957998

0xf7593d26

0xf7712d60

0x0804868c

0xff957978

//Buffer overflow

0x61593d00

// "a"

0x61616161

// "aaaa"

0x61616161

// "aaaa"

0x61616161

// "aaaa"

0x0804847b

// "\x7b\x84\x04\x08", *evil_code* address

Argh, we got hacked!

//Success! *evil_code* was executed

Segmentation fault (core dumped)

Modern ROP attacks

- ▶ ROP attacks are a common kind of attack in the industry
- ▶ Modern ROP attacks are much more complicated
- ▶ *Return-to-libc* attacks uses code from the *glibc* library to construct malicious code and uses return addresses to execute it


Goals of our approach

We want to have a way to protect return addresses at runtime.

- ▶ Modifications of the memory layout in order to have an easy way to know return addresses location
- ▶ Code transformations which add runtime checks on the dangerous instruction in order to forbid any illegal write on the return addresses locations

CompCert stack

- ▶ Programs memory is separated into multiple area like the heap, the stack or the code section
- ▶ Return addresses are solely located in the stack
- ▶ The stack is composed of piled up frames each related to a function being executed
- ▶ Frames store data of their respective function



Transformations of the stack layout

All the return addresses locations a verify the equality $a \bmod n = 0$, with n a constant offset between the return addresses locations.

1. Set a constant offset n between all the return addresses
2. Align the stack

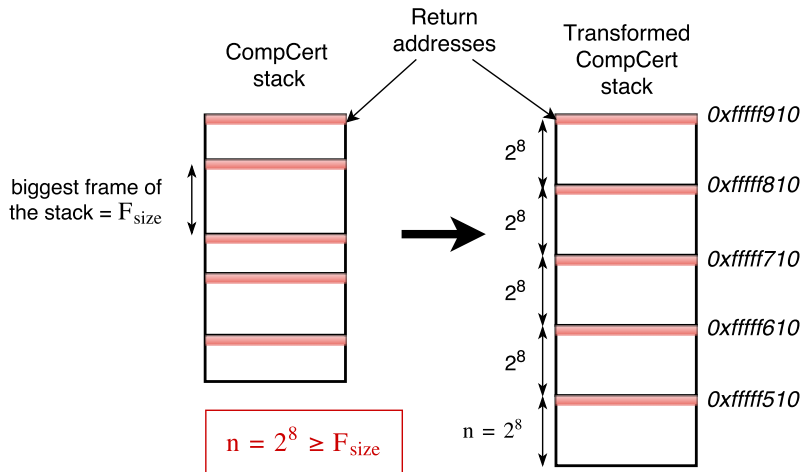
Constant offset n between return addresses (1/2)

Constant offset n between return addresses locations

- ▶ Fix frames size to n
- ▶ Pick n as the biggest frame size of the previous stack
- ▶ Pick n as a power of two

With this we have $a \bmod n = c$, with c the location of the first return address in the stack

Constant offset n between return addresses (2/2)

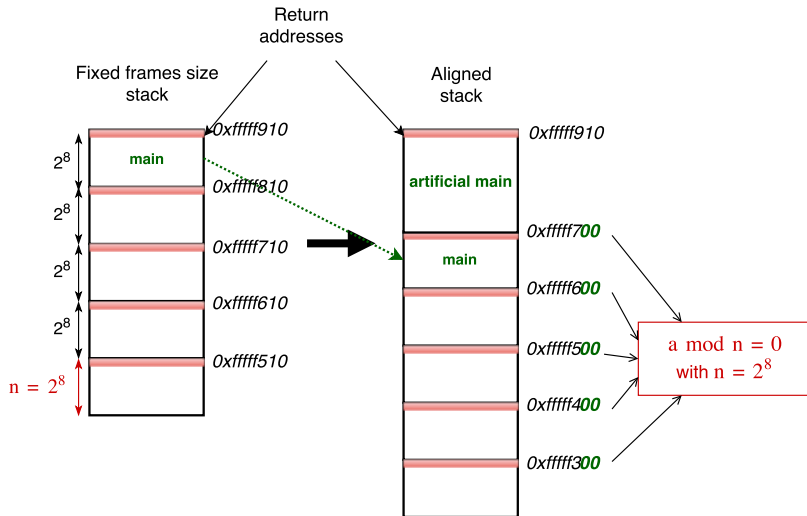


Stack alignment (1/2)

We currently have the equality $a \bmod n = c$ but we want $a \bmod n = 0$ with a any return address locations and c the first return address location in the stack.

- ▶ introduce an artificial function at the beginning of the program
- ▶ the function align the stack as we wanted
- ▶ the function then calls the *main* of the program

Stack alignment (2/2)



Injection of runtime checks

1. Check if the address is part of the stack
2. Check if the address verifies $a \bmod n = 0$

```
1  if (targeted_address > 0xff000000) {  
2      temp_var = targeted_address & (n-1);  
3      if (temp_var < 3) {  
4          Error behaviour  
5      }  
6  }  
7  *targeted_address = value;  
8  Continue execution ...
```

Branchless runtime checks

In certain cases branchless code shows much better performance

```
1  if (targeted_address > 0xff000000) {  
2      temp_var = targeted_address & (n-1);  
3      temp_var = temp_var - 3;  
4      temp_var = temp_var >> 31;  
5      temp_var = ~temp_var;  
6      targeted_address = temp_var &  
          targeted_address;  
7  }  
8  *targeted_address = value;  
9  Continue execution ...
```


Conditions of our approach

- ▶ No modifications of the stack (inline assembly)

```
1 int foo(int a) {  
2     asm( '\ $sub 50, \ %esp ' );  
3     printf( " Hello world!" );  
4 }
```

- ▶ Need to recompile extern libraries with the same frames size

Evaluation of performance (1/2)



Discussion

- ▶ Test our implementation against more complicated ROP attacks
- ▶ Reduce the number of runtime checks with static analysis
- ▶ Improve the performance of the runtime checks with a super-optimizer
- ▶ See the impact of our approach on memory consumption